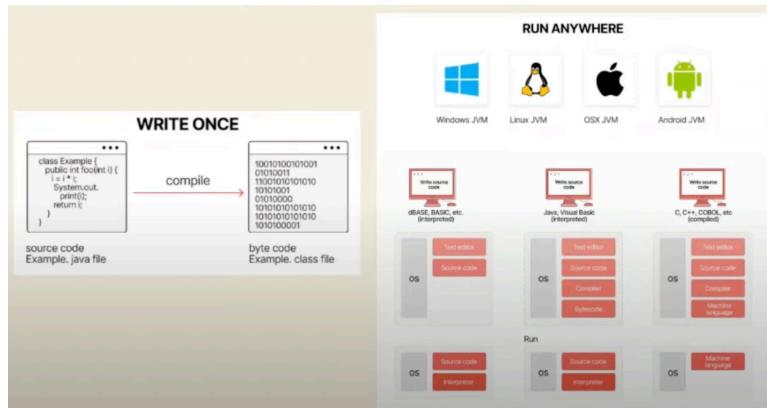


The Role of Ethereum Founders

In the domain of blockchain technology, the names Vitalik Buterin and Gavin Wood are of significant importance. While Buterin's role as the co-founder of Ethereum enjoys widespread recognition, Wood's contributions often remain in the background. But, Gavin Wood had a pivotal role in shaping the Ethereum ecosystem. In 2013, when Vitalik Buterin introduced Ethereum at the age of 19, he was not alone in this endeavor. Gavin Wood stood alongside him as a foundational figure whose significance parallels that of Buterin. Wood's instrumental role in the design of the Ethereum Virtual Machine (EVM) breathed life into Ethereum's core infrastructure. The EVM sustains Ethereum's functionality, rendering it indispensable to the network's operations. Furthermore, Wood's contribution extended beyond the EVM: he is credited as the architect behind the Solidity language, the programming language essential for crafting smart contracts on the Ethereum platform.

The History of Virtual Machines



A virtual machine serves as a program emulating the functionality of a computer by encapsulating its genuine components and operating system. It operates as an autonomous entity atop the host computer, irrespective of its underlying operating system. For instance, a physical computer equipped with an OS like Windows can concurrently host multiple virtual machines, each running with distinct operating systems. This independence underscores the versatility and scalability of virtual machine technology. Virtual machines exhibit diversity in their capabilities, tailored to fulfill specific tasks and functionalities. For instance, VirtualBox, a prominent example, facilitates the creation of virtual machines atop hardware, enabling the execution of various operating systems. Similarly, Java Virtual Machines provide a runtime environment for Java programs, facilitating cross-platform execution.

Every virtual machine possesses inherent commonalities. They operate autonomously, existing in isolation from other computational components. They exclusively leverage computational resources, diverging significantly from the execution paradigm of native programs that directly interact with underlying hardware components through the OS. For instance, consider a scenario where an operating system hosts an application. In this conventional setup, the application operates within the border of the host system, utilizing its computational resources. However, in a virtualized environment, the virtual machine encapsulates its computational resources, independent of the underlying host system. Consequently, applications running atop a virtual machine leverage the computational resources inherent to the virtual environment rather than those of the host system, thus seeing the virtual machine as a regular process.

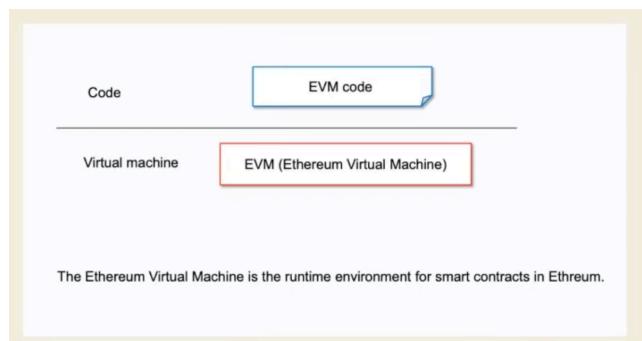
The compilation of software is also different. Traditional development entails translating source code into platform-specific machine code, thereby restricting portability across diverse architectures. Conversely, within the realm of virtual machines, source code undergoes compilation into bytecode, serving as an intermediary representation. This bytecode abstraction facilitates platform independence, as it can be executed on any platform leveraging the virtual machine environment. Consequently, the onus of generating platform-specific machine code for

The Ethereum Protocol and the Ethereum Virtual Machine

diverse architectures is alleviated, underscoring the pivotal role of bytecode as a universal intermediary.

The significance of bytecode transcends mere abstraction, manifesting as a cornerstone for creating an isolated and universally compatible execution environment. By circumventing platform-specific constraints, bytecode streamlines the development process, enabling developers to concentrate on code creation without concerns regarding underlying hardware architectures or operating system idiosyncrasies.

The Ethereum Virtual Machine



In simplified terms, the Ethereum Virtual Machine (EVM) serves as a runtime environment for executing smart contracts within the Ethereum ecosystem. Drawing parallels with the Java Virtual Machine (JVM), which facilitates the execution of Java programs across different platforms, the EVM enables the execution of EVM bytecode on various machines. This fundamental function of the EVM underscores its role as an isolated entity within the system, ensuring that operations conducted within the EVM do not impact data or programs elsewhere.

The isolation of the EVM signifies its independence from other system components, akin to a standalone machine or environment. Consequently, operations within the EVM remain insulated, regardless of the frequency or nature of function calls. The EVM functions as a computational entity facilitating data storage on the blockchain, emphasizing its pivotal role within the Ethereum ecosystem.

In practice, developers compose smart contracts in languages like Solidity or Vyper, which are then compiled into bytecode by the EVM. This bytecode, comprising a collection of opcodes, directs the EVM's operations, influencing changes in the Ethereum state. While bytecode execution is integral to the EVM's functionality, its platform-independent nature ensures uniformity across diverse machines, mitigating concerns regarding underlying CPU architectures. Within the Ethereum ecosystem, the EVM operates as a distributed entity, orchestrating consensus among Ethereum nodes to validate changes in the blockchain state.

The execution of operations on the EVM necessitates the provision of gas, representing computational costs within the Ethereum network. In practical terms, the integration of the EVM into Ethereum client services like Geth facilitates its deployment and operation within the network. This symbiotic relationship underscores the EVM's significance within Ethereum's decentralized infrastructure.

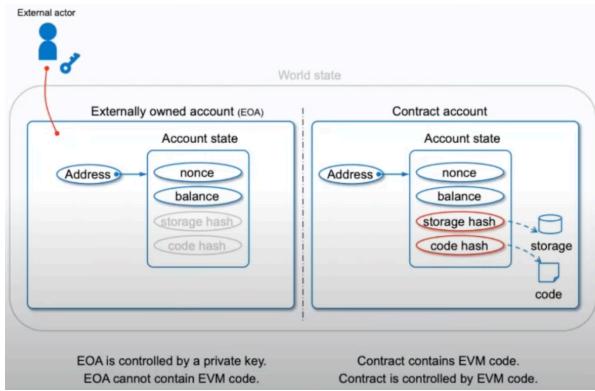
Ethereum and the Account-based Ledger

Each account, as delineated, encompasses several key attributes, including address, nonce, balance, storage hash, and code hash. However, a notable disparity emerges between externally owned accounts and contract accounts.

Externally owned accounts epitomize user-controlled entities devoid of associated code or storage hash. These accounts, typified by those generated through platforms like MetaMask, are

The Ethereum Protocol and the Ethereum Virtual Machine

governed by individual private keys, granting exclusive authority over account management. Conversely, *Contract Accounts*, while sharing common attributes with externally owned counterparts, diverge notably through the inclusion of code and storage hash. Contract accounts, established upon deployment of smart contracts, encapsulate executable code and associated storage, administered not by private keys but by the logic encoded within the contract.

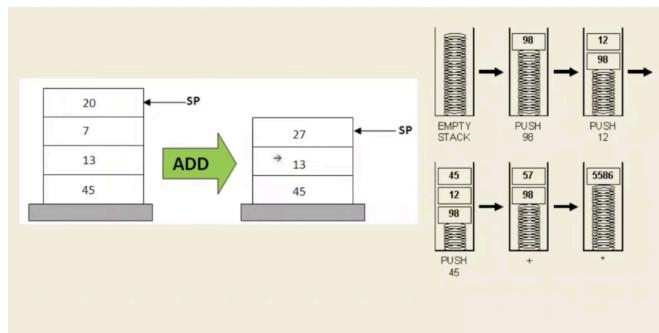


This fundamental contrast underscores the distinct governance mechanisms inherent in each account type. Whereas externally owned accounts rely on private key control, contract accounts defer to the instructions delineated within their respective codebases. This functional dichotomy resonates with the broader ethos of blockchain technology, where code functions as the arbiter of contractual agreements and operational parameters.

Transacting with externally owned accounts entails adherence to stringent signing protocols, necessitating cryptographic validation through private keys. Conversely, interactions with contract accounts involve invocation of predefined functions dictated by the embedded codebase, bypassing conventional private key validation procedures.

Furthermore, contract accounts assume multifaceted roles within the Ethereum ecosystem, serving not merely as repositories for cryptocurrency balances but as conduits for deploying and executing smart contracts. This dynamic functionality underscores the pivotal role of contract accounts in facilitating decentralized application (dApp) development and execution.

EVM Foundations: Stack Based Machine



Two primary paradigms characterize virtual machines: *register-based* and *stack-based* architectures.

Register-based virtual machines utilize a designated set of registers for variable storage and computation. Predominantly associated with Intel architecture, these virtual machines employ registers as primary storage units for operand manipulation and calculation. This architecture

The Ethereum Protocol and the Ethereum Virtual Machine

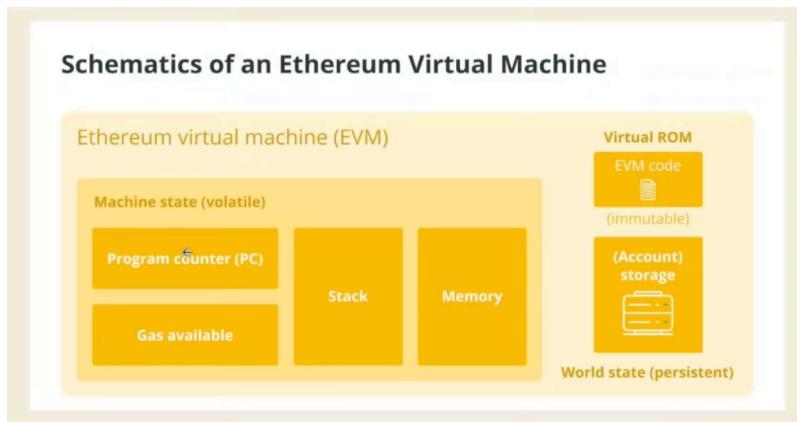
harnesses the efficiency of direct register access, facilitating rapid computation and data manipulation.

Conversely, *stack-based* virtual machines rely on a stack data structure for variable storage and computation. In this model, data is stored and retrieved using Last In, First Out (LIFO) semantics, wherein the last item inserted into the stack is the first to be removed. Notable examples of stack-based machines include the Java Virtual Machine (JVM), renowned for its utilization of stack-based execution mechanisms.

A distinguishing characteristic of stack-based architectures is the perpetual availability of data at the top of the stack. This inherent property simplifies operand retrieval and facilitates streamlined computation, as operands crucial to arithmetic operations consistently reside at the stack's apex. The stack pointer (SP) governs stack navigation, consistently referencing the topmost element within the stack structure.

The Ethereum Virtual Machine (EVM) exemplifies a stack-based architecture, embodying the principles of stack-based computation. Designed as a simple stack-based virtual machine, the EVM leverages stack-based execution mechanisms to facilitate efficient computation and state manipulation within the Ethereum ecosystem.

EVM Architecture



The Ethereum Virtual Machine (EVM) comprises several integral components, each playing a crucial role in the execution and management of smart contracts and decentralized applications (dApps).

Program Counter (PC): At the heart of the EVM lies the program counter, a register that points to the current instruction or opcode to be executed within the bytecode of a smart contract. As the contract runs, the program counter increments, ensuring sequential execution of instructions.

Gas Available: Gas is a fundamental concept in Ethereum, representing the computational resources allocated for executing transactions and smart contracts. Gas available denotes the amount of gas provided for a specific transaction, determining its execution capability within the EVM.

Stack: The EVM employs a stack data structure for temporary storage and operand manipulation during smart contract execution. Utilizing a Last In, First Out (LIFO) approach, the stack facilitates efficient storage and retrieval of data, aiding in computation and output generation. Each call context in the EVM spawns a distinct stack instance, ensuring isolation and memory management efficiency.

Memory: EVM memory serves as temporary, ephemeral storage for computational operations and data manipulation during contract execution. Similar to the stack, memory is ephemeral and is

The Ethereum Protocol and the Ethereum Virtual Machine

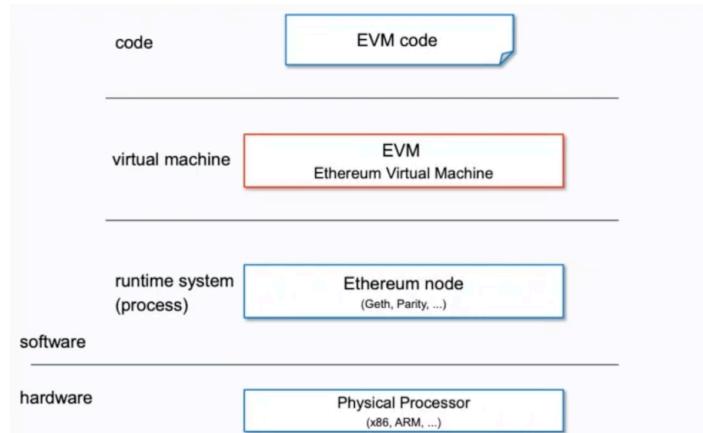
cleared at the conclusion of each call context, ensuring resource optimization and memory integrity.

Storage: Persistent memory is essential for maintaining the state of smart contracts across successive transactions. Storage within the EVM represents the long-term, persistent data repository for individual contracts, allowing for the retention and retrieval of critical contract data and state information. Each smart contract possesses its unique storage space, ensuring data integrity and contract autonomy.

Code: The bytecode of smart contracts resides within the EVM's code segment, comprising the executable instructions that define contract behavior and functionality. As part of the contract account state, the code segment is persistent and immutable, housing the essential logic and operations encoded within smart contracts. Bytecode execution within the EVM dictates contract behavior and facilitates transaction processing and state updates.

Each component of the EVM architecture contributes to the seamless execution and management of smart contracts within the Ethereum ecosystem. From instruction processing to resource allocation and state management, these components collectively underpin the functionality and efficiency of decentralized applications and smart contract execution on the Ethereum blockchain.

EVM Role in the Ethereum Architecture



EVM Code: The EVM code comprises the executable instructions and logic that define the behavior of the Ethereum Virtual Machine. This code is an integral part of Ethereum's protocol and is responsible for executing smart contracts and processing transactions.

Ethereum Virtual Machine: The EVM functions as a virtual CPU or a runtime environment within Ethereum nodes. It interprets and executes the bytecode of smart contracts, facilitating the decentralized execution of code across the Ethereum network.

Runtime Processes (Ethereum Nodes): Ethereum nodes, such as Geth, Parity, and other client services, host the EVM as a runtime process. These nodes maintain a synchronized copy of the Ethereum blockchain and actively participate in transaction validation, block propagation, and consensus mechanisms.

Hardware Infrastructure: Each Ethereum node runs on physical or virtual hardware infrastructure, comprising processors, memory, storage, and network connectivity. These resources provide the necessary computational power and storage capacity to host the EVM and support the execution of smart contracts and transaction processing.

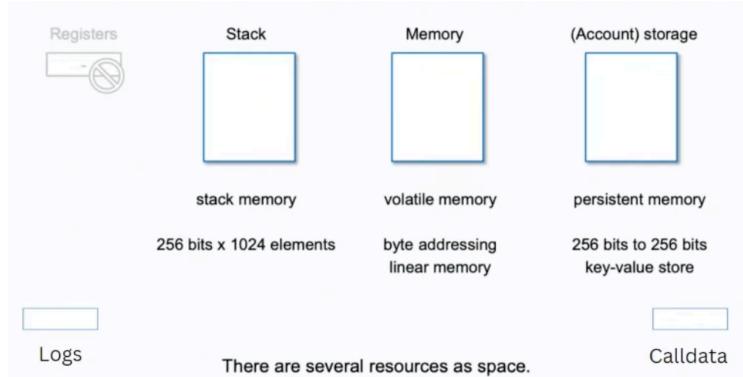
In essence, the EVM operates as a decentralized computing engine embedded within Ethereum nodes, facilitating the execution of smart contracts and the processing of transactions across the network. As part of the broader Ethereum ecosystem, the EVM interacts with runtime processes,

The Ethereum Protocol and the Ethereum Virtual Machine

client services, and underlying hardware infrastructure to enable the seamless operation of decentralized applications and services on the Ethereum blockchain.

Solidity Code: the programming language used to write smart contracts, is compiled into bytecode compatible with the EVM. This bytecode is then distributed across the Ethereum network to be executed by individual Ethereum nodes running the EVM. Through this distributed execution model, the EVM ensures consensus and integrity in smart contract execution while leveraging the computational resources of the decentralized Ethereum network.

EVM Storage



How the Virtual Machine (EVM) access and store information during runtime? Let's summarize and elaborate on each of these components:

Stack: Acting as a temporary storage mechanism, the stack operates on a last-in-first-out (LIFO) basis. It stores values temporarily during runtime and serves as an intermediate for reading, writing, and manipulating data from other spaces like memory, storage, call data, and logs.

Memory: Similar to RAM in traditional computer systems, memory provides temporary storage during runtime. It is volatile and linearly addressed, allowing for the storage and retrieval of data during the execution of smart contracts.

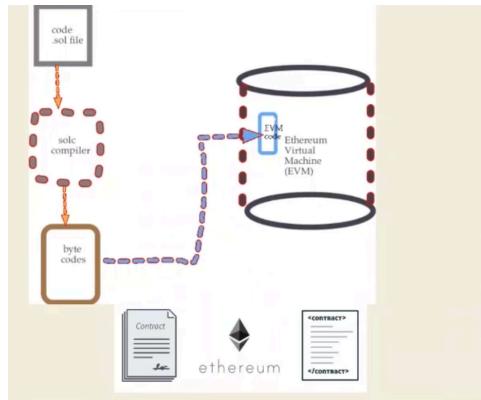
Call Data: Call data refers to the data field of a transaction and stores the input parameters or arguments of a function to be called. It is read-only and cleared after the execution of the transaction, acting as a temporary memory for function arguments.

Logs: Logs are write-only data areas used to emit log events from smart contracts. These events are often specified in Solidity code using the `emit` keyword and can be used for debugging, auditing, or interfacing with front-end applications.

Storage: Storage serves as persistent memory for smart contracts, allowing them to store data permanently between transactions. It operates as a key-value store and is represented in the form of a Merkle Patricia Trie. While the blockchain itself stores only block headers, clients like Geth and Aragon store the full trie contents in a database, ensuring efficient data retrieval and manipulation.

Understanding these machine spaces is crucial for developers working with smart contracts on the Ethereum blockchain. It enables them to optimize their code, manage data efficiently, and leverage the capabilities of the EVM effectively. Additionally, familiarity with these concepts allows developers to navigate the complexities of blockchain development and build robust decentralized applications.

Solidity



Solidity is indeed a high-level, object-oriented programming language specifically designed for writing smart contracts on the Ethereum blockchain and other EVM-compatible chains. Here's a breakdown of some key points about Solidity:

Purpose: Solidity is primarily used for creating smart contracts, which are self-executing contracts with the terms of the agreement between buyer and seller directly written into code. These contracts automatically execute and enforce themselves when predetermined conditions are met.

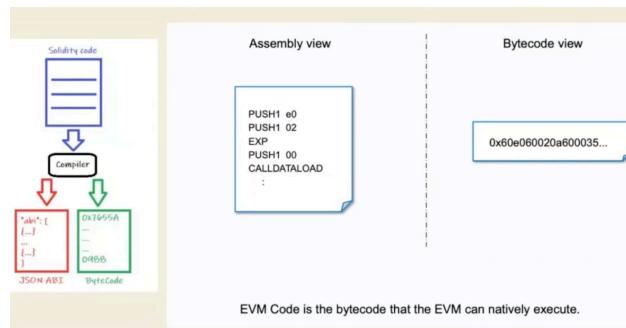
Object-Oriented: Solidity supports many features of object-oriented programming, including inheritance, polymorphism, and encapsulation. This makes it easier for developers to structure their code and manage complex relationships between different parts of the contract.

Similarities to Other Languages: Solidity syntax is similar to that of other programming languages like C++, Java, and JavaScript. This familiarity makes it easier for developers with experience in these languages to transition to Solidity.

Compilation: Solidity code is compiled into bytecode using the Solidity compiler. This bytecode is then deployed onto the Ethereum Virtual Machine (EVM) or other EVM-compatible blockchains. The EVM executes the bytecode, enabling the smart contract to perform its predefined functions.

Integration with Ethereum: Solidity is the most widely used language for writing smart contracts on the Ethereum blockchain. It provides developers with the tools and capabilities needed to create decentralized applications (DApps), token contracts, governance systems, and more.

The Bytecode



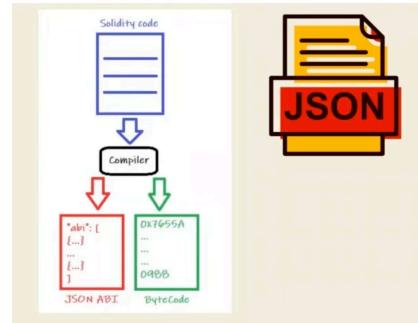
Let us now delve into the concept of bytecode, a fundamental aspect of Ethereum smart contracts. Contrary to conventional programming paradigms where source code is directly deployed, Ethereum smart contracts are instantiated through bytecode on the Ethereum blockchain. Bytecode, represented in hexadecimal notation, embodies the compiled form of the smart contract codebase.

The Ethereum Protocol and the Ethereum Virtual Machine

Upon compilation of Solidity smart contracts, the source code undergoes translation into bytecode, which is exclusively interpretable by the Ethereum Virtual Machine (EVM). It is imperative to note that EVM bytecode is not directly executable by conventional processors; instead, it necessitates an intermediary program. This program can either interpret the bytecode directly within the virtual machine or further compile it into machine code for enhanced performance.

The compilation process typically yields bytecode alongside an Application Binary Interface (ABI), which facilitates interaction with the smart contract. The bytecode, as depicted in the example, can be observed alongside its corresponding assembly view, providing insights into the breakdown of opcodes within the bytecode. This decomposition of bytecode into constituent opcodes enables developers to comprehensively analyze and optimize contract logic.

The Application Binary Interface (ABI)



ABI serves as an essential interface facilitating interaction with smart contracts post-compilation of Solidity code. Upon compilation, Solidity code yields two crucial components: a JSON API and bytecode. The ABI, an acronym for Application Binary Interface, acts as a conduit for interacting with smart contracts. It encapsulates essential information pertaining to contract functions, variables, modifiers, and parameters in a structured JSON format. This JSON representation serves as a comprehensive descriptor facilitating seamless communication between high-level web applications, typically scripted in languages such as JavaScript or TypeScript, and the bytecode executed within the Ethereum Virtual Machine (EVM). Essentially, the ABI bridges the semantic gap between web applications and EVM bytecode, enabling effective communication and interaction. Thus, ABI serves as a pivotal intermediary, fostering interoperability and facilitating the seamless integration of decentralized applications (dApps) with Ethereum smart contracts.

The OpCode

An interactive reference to Ethereum Virtual Machine OpCodes							
opcode	name	gas cost	stack depth	memory	storage	description	example
0x00	STOP	0	0	0	0	Halts execution	
0x01	ADD	2	1 → 1	0	0	Addition operation	
0x02	MUL	5	1 → 1	0	0	Multiplication operation	
0x03	SUB	3	1 → 1	0	0	Subtraction operation	
0x04	DIV	5	1 → 1	0	0	Integer division operation	
0x05	SDIV	5	1 → 1	0	0	Signed integer division operation (truncated)	
0x06	MOD	5	1 → 1	0	0	Modulo remainder operation	
0x07	SMOD	5	1 → 1	0	0	Signed modulo remainder operation	
0x08	ADDMOD	8	1 → 1 → 1	0	0	Modulo addition operation	
0x09	MULMOD	8	1 → 1 → 1	0	0	Modulo multiplication operation	
0x0A	EXP	10 ⋯	1 → exponent	0	0	Exponentiation operation	
0x0B	SIGEXTEND	5	1 → 1	0	0	Extend length of left complement signed	

Let us now explore the concept of opcode, a fundamental aspect within the realm of Ethereum bytecode. The term "opcode" is an abbreviation for "operation code," denoting a type of code instructing the machine on specific actions to execute. Ethereum bytecode, being no exception,

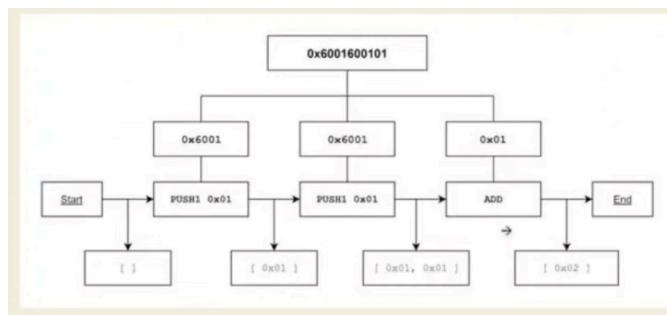
The Ethereum Protocol and the Ethereum Virtual Machine

can be dissected into a series of operands and opcodes, each serving distinct operational functions.

The realm of opcodes encompasses various categories, including stack management, arithmetic operations, environmental directives, memory management, and more. Given the breadth of opcode types and functionalities, a comprehensive understanding warrants dedicated study, potentially meriting a specialized crash course.

For elucidation, let us consider a byte code example. Each opcode within the bytecode possesses a designated mnemonic representation, alongside its associated gas cost. For instance, the opcode "0x01" corresponds to the mnemonic "add," signifying addition operation within the Ethereum Virtual Machine (EVM). Reference to resources such as the Ethereum website furnishes comprehensive lists detailing opcode functionalities and their respective mnemonics.

In essence, opcodes encapsulate the core instructions for performing operations within the EVM bytecode, exemplifying the intricate interplay between bytecode and opcode representations. Should a deeper exploration into opcode intricacies be desired, a dedicated crash course could be formulated to delve into the nuanced facets of opcode functionality within Ethereum bytecode.



Let us delve into the operational dynamics of specific opcodes within Ethereum bytecode, elucidating their pivotal role in stack manipulation and gas consumption optimization. One such opcode of significance is "swap," which performs the function of removing two elements from the stack - namely, the last and first elements - and subsequently replaces them with the result. This mechanism ensures that the resultant value occupies the topmost position on the stack, adhering to the stack's LIFO (Last In, First Out) structure.

The intricacies of EVM bytecode, particularly in terms of gas consumption minimization, merit thorough comprehension. Mastery of opcode functionalities facilitates the crafting of bytecode sequences optimized for minimal gas expenditure, thereby enhancing cost-efficiency for end users. Notably, opcodes such as "push" exemplify this optimization principle, wherein the insertion of data onto the stack is accomplished efficiently.

Furthermore, it is imperative to grasp the overarching objective of gas optimization, which extends beyond merely reducing deployment costs. The paramount aim is to mitigate gas expenses incurred by end users during smart contract interactions. A judiciously optimized contract not only minimizes deployment gas but also ensures a favorable user experience by mitigating transactional costs.

In essence, the symbiotic relationship between bytecode and opcodes underscores the imperative of crafting optimally efficient source code. By meticulously optimizing Solidity code, developers not only alleviate gas burdens for end users but also foster broader adoption and interaction with their smart contracts. Thus, the optimization endeavor transcends mere cost reduction, constituting a pivotal aspect of facilitating seamless and cost-effective blockchain interactions.

The Assembly Language

```
101     function _name() internal pure override returns (string memory) {
102         // Return the name of the contract.
103         assembly {
104             mstore(0x20, 0x20)
105             mstore(0x47, 0x07536561706f7274)
106             return(0x20, 0x60)
107         }
108     }
```

Another aspect to consider is the significance of Assembly Language in the context of smart contract development. Assembly Language bears importance due to its inherent portability, facilitating the authoring of entire smart contracts utilizing optimized opcode sequences (OPCODES). This underscores the pivotal role of Solidity Assembly Language, albeit its complexity, in augmenting smart contract functionality beyond the capabilities of Solidity alone. Moreover, employing Assembly Language can yield efficiency gains in terms of gas consumption, a critical consideration in blockchain environments.

It is evident that the utilization of Assembly Language, commonly referred to as Yul within the Solidity ecosystem, is integral to advanced contract development. Notably, Yul code blocks are encapsulated within curly braces, akin to other programming constructs, facilitating its integration within Solidity contracts. For comprehensive insights into Yul and its syntax, reference to the Solidity documentation is recommended. The prescribed format for integrating Yul within Solidity contracts involves encapsulating the Yul code within an "assembly" block delineated by curly braces, thereby fostering seamless integration of optimized opcode sequences within the contract logic.